

Algorithm 2025 Spring reference answers

1.

LCS = 001011

2.

We assume that the actual cost of a search equals the number of nodes examined, i.e., the depth of the node found by the search in T , plus 1. Then the expected cost of a

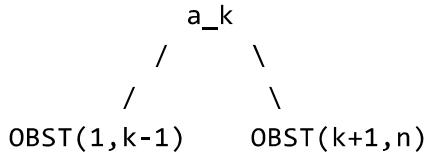
$$\text{search in } T \text{ is } E[C_T] = \sum_{i=1}^k (l_i + 1) \cdot p_i = \sum_{i=1}^k p_i + \sum_{i=1}^k l_i \cdot p_i = 1 + \sum_{i=1}^k l_i \cdot p_i$$

And the recurrence of OBST:

- $e(i, j) = \begin{cases} p_i, & i = j \\ \min_{1 \leq k \leq n} \{e(i, k - 1) + e(k + 1, j) + w(i, j)\}, & i < j \end{cases}$
- $w(i, j) = \sum_{k=i}^j p_k$

(a) (answer 6%, explain 4%)

- In an optimal binary search tree on the ordered keys a_1, a_2, \dots, a_n , we can denote it by $\text{OBST}(1, n)$
- If we decide that a_k is a root, then the entire tree decomposes into two independent subtrees:
 1. Left subtree
 - Contains the keys a_1, a_2, \dots, a_{k-1} .
 - It must be the OBST for those keys
 - We can denote it by $\text{OBST}(1, k - 1)$
 2. Right subtree
 - Similarly, we can denote it by $\text{OBST}(k + 1, n)$
- Graphically, the structure is:



(b) (answer 6%, explain 4%)

Suppose in the OBST we have chosen a_k as the root, then we can denote

- $W_L = \sum_{a_i \in L} p_i = \sum_{i=1}^{k-1} p_i$
- $W_R = \sum_{a_i \in R} p_i = \sum_{i=k+1}^n p_i$

Let p_k be the weight of the root key a_k . Then when we splice L and R under a_k , every node in L and R moves "**one level deeper**", so each of them incurs an extra p_k in the total cost. Meanwhile a_k sits at the depth 1 and contributes p_k . Hence

- $C_T = (C_L + W_L) + p_k + (C_R + W_R)$

But note that

- $W_T = W_L + p_k + W_R$

so we can also write

- $C_T = C_L + C_R + W_T$

According to the recurrence of OBST, We can say

- $C_{OBST} = e(1, n)$
- $e(1, k - 1)$ plays the role of C_L
- $e(k + 1, n)$ plays the role of C_R
- $w(1, n)$ is exactly W_T

(C) (answer 6%, explain 4%)

When all weights are equally likely, $p_i = \frac{1}{n}$, the total weight of the tree on n keys is

- $W_T = \sum_{i=1}^n p_i = n \cdot \frac{1}{n} = 1$

Recall that if you choose a_k as the root, then

- $C_T = C_L + C_R + W_T = C_L + C_R + 1$

Since all the weights are the same, we can define

- $C_{OBST} = C(n) = \begin{cases} 0, & n = 0 \\ \min_{1 \leq k \leq n} \{C(k-1) + C(n-k) + 1\}, & n > 0 \end{cases}$

Because the optimal split is always as balanced as possible when all p_i are the same, so we get exactly the recurrence:

- $C(n) = \begin{cases} 0, & n = 0 \\ C(\lfloor \frac{n-1}{2} \rfloor) + C(\lceil \frac{n-1}{2} \rceil) + 1, & n > 0 \end{cases}$

After solving the recurrence,

- $C(n) = \Theta(\log n)$

3.

Yes, if we sort the nodes with $O(n)$ sorting algorithms, we can build the balance tree with a $T(n) = 2T(n/2) + O(1) = O(n)$ algorithm:

```
build_tree(sorted_list):
    N = sorted_list.size()

    if N == 1:
        root.val = sorted_list[0]
        root.left = root.right = null
        return root
    if N <= 0:
        return null

    idx = floor(N/2)
    root.val = sorted_list[idx]
    root.left = build_tree(sorted_list[:idx])
    root.right = build_tree(sorted_list[idx+1:])
    return root
```

Or even better, trivial $O(n)$ algorithm to build a crooked tree:

```

build_tree(sorted_list):
    if sorted_list.size() == 0:
        return null
    root.val = sorted_list[0]
    root.left = null
    root.right = build_tree(sorted_list[1:])
    return root

```

4.

Algorithm Steps (6%)

1. Decrease x to $-\infty$
 - Set key[x] to a value smaller than every other key.
 - Then "bubble" x up to the root list.
2. Extract the minimum.
 - Now x is the minimum element in H , and it removes x from the root list.
 - Take x 's children and merges them back into H .

Complexity Analysis (4%)

Let n be the size of the heap before deletion.

1. Decrease-key to $-\infty$:
 - A binomial tree of order k has height k (property 2).
 - There are at most $\lfloor \log_2 n \rfloor$ orders, i.e. $k = O(\log n)$, in the heap.
 - So the cost is $O(\log n)$.
2. Extract-min:

Since there are at most $O(\log n)$ roots the root list, in this operation we need:

 - Scanning the root list of H .
 - Splice out and reversing the list of x 's children, $O(\log k)$.
 - Merging two root-lists (original roots and children).
 - And the "linking" steps are constant time.
 - So the total cost is $O(\log n)$

5.

(a)

GENERIC-MST(G, w)

- 1 $A = \emptyset$
- 2 **while** A does not form a spanning tree
- 3 find an edge (u, v) that is safe for A
- 4 $A = A \cup \{(u, v)\}$
- 5 **return** A

Theorem 23.1

Let $G = (V, E)$ be a connected, undirected graph with a real-valued weight function w defined on E . Let A be a subset of E that is included in some minimum spanning tree for G , let $(S, V - S)$ be any cut of G that respects A , and let (u, v) be a light edge crossing $(S, V - S)$. Then, edge (u, v) is safe for A .

Proof Let T be a minimum spanning tree that includes A , and assume that T does not contain the light edge (u, v) , since if it does, we are done. We shall construct another minimum spanning tree T' that includes $A \cup \{(u, v)\}$ by using a cut-and-paste technique, thereby showing that (u, v) is a safe edge for A .

The edge (u, v) forms a cycle with the edges on the simple path p from u to v in T , as Figure 23.3 illustrates. Since u and v are on opposite sides of the cut $(S, V - S)$, at least one edge in T lies on the simple path p and also crosses the cut. Let (x, y) be any such edge. The edge (x, y) is not in A , because the cut respects A . Since (x, y) is on the unique simple path from u to v in T , removing (x, y) breaks T into two components. Adding (u, v) reconnects them to form a new spanning tree $T' = T - \{(x, y)\} \cup \{(u, v)\}$.

We next show that T' is a minimum spanning tree. Since (u, v) is a light edge crossing $(S, V - S)$ and (x, y) also crosses this cut, $w(u, v) \leq w(x, y)$. Therefore,

$$\begin{aligned} w(T') &= w(T) - w(x, y) + w(u, v) \\ &\leq w(T). \end{aligned}$$

(b)

```
function Kruskal(Graph G) is
    F := ∅
    for each v in G.Vertices do
        MAKE-SET(v)

        for each {u, v} in G.Edges ordered by increasing weight({u, v}) do
            if FIND-SET(u) ≠ FIND-SET(v) then
                F := F ∪ { {u, v} }
                UNION(FIND-SET(u), FIND-SET(v))
    return F
```

Find minimum [edit]

To find the **minimum** element of the heap, find the minimum among the roots of the binomial trees. This can be done in $O(\log n)$ time, as there are just $O(\log n)$ tree roots to examine.^[1]

By using a pointer to the binomial tree that contains the minimum element, the time for this operation can be reduced to $O(1)$. The pointer must be updated when performing any operation other than finding the minimum. This can be done in $O(\log n)$ time per update, without raising the overall asymptotic running time of any operation.[citation needed]

Delete minimum [edit]

To **delete the minimum element** from the heap, first find this element, remove it from the root of its binomial tree, and obtain a list of its child subtrees (which are each themselves binomial trees, of distinct orders). Transform this list of subtrees into a separate binomial heap by reordering them from smallest to largest order. Then merge this heap with the original heap. Since each root has at most $\log_2 n$ children, creating this new heap takes time $O(\log n)$. Merging heaps takes time $O(\log n)$, so the entire delete minimum operation takes time $O(\log n)$.^[1]

```
function deleteMin(heap)
    min = heap.trees().first()
    for each current in heap.trees()
        if current.root < min.root then min = current
    for each tree in min.subTrees()
        tmp.addTree(tree)
    heap.removeTree(min)
    merge(heap, tmp)
```

- Time Complexity: $O(E \log E)$

(c)

```
1  function Prim(vertices, edges) is
2      for each vertex in vertices do
3          cheapestCost[vertex] ← ∞
4          cheapestEdge[vertex] ← null
5
6      explored ← empty set
7      unexplored ← set containing all vertices
8
9      startVertex ← any element of vertices
10     cheapestCost[startVertex] ← 0
11
12     while unexplored is not empty do
13         // Select vertex in unexplored with minimum cost
14         currentVertex ← vertex in unexplored with minimum cheapestCost[vertex]
15         unexplored.remove(currentVertex)
16         explored.add(currentVertex)
17
18         for each edge (currentVertex, neighbor) in edges do
19             if neighbor in unexplored and weight(currentVertex, neighbor) < cheap
20                 cheapestCost[neighbor] ← weight(currentVertex, neighbor)
21                 cheapestEdge[neighbor] ← (currentVertex, neighbor)
22
23         resultEdges ← empty list
24         for each vertex in vertices do
25             if cheapestEdge[vertex] ≠ null THEN
26                 resultEdges.append(cheapestEdge[vertex])
27
28     return resultEdges
```

- Maintain cheapestCost with binomial heap

Decrease key [\[edit\]](#)

After **decreasing** the key of an element, it may become smaller than the key of its parent, violating the minimum-heap property. If this is the case, exchange the element with its parent, and possibly also with its grandparent, and so on, until the minimum-heap property is no longer violated. Each binomial tree has height at most $\log_2 n$, so this takes $O(\log n)$ time.^[1] However, this operation requires that the representation of the tree include pointers from each node to its parent in the tree, somewhat complicating the implementation of other operations.^[3]

- Time Complexity: $O(V \log V) + O(E \log V) = O(E \log V)$

6.

Algorithm setup

- We maintain a parent pointer $p[x]$ for each node x .
 - If $p[x] = x$, x is a root.

- Otherwise, $p[x]$ points towards its parent in the tree.
- For each root r we keep $\text{size}[r] =$ the number of nodes in r 's tree.
- To union two trees rooted at r_1, r_2 .
 1. Compare sizes: assume $\text{size}[r_1] \leq \text{size}[r_2]$.
 2. Make r_1 's parent point to r_2 : $p[r_1] \leftarrow r_2$
 3. Update $\text{size}[r_2] \leftarrow \text{size}[r_1] + \text{size}[r_2]$

Claim: Total cost of $n-1$ unions is $O(n)$ (3%)

- Starting from n singleton trees, you need exactly $n-1$ unions to merge them all into one tree. Since each union is just one constant-time pointer assignment ($p[r_1] \leftarrow r_2$), the total work is $(n - 1) \times O(1) = O(n)$

Claim: Height bound: $h \leq \lfloor \log n \rfloor$ (7%)

- <https://hackmd.io/@KentLee/Syshq-trJx> (<https://hackmd.io/@KentLee/Syshq-trJx>).

7.

When k is a constant, the problem is easy; otherwise, it is difficult since if it's easy, we can make $k = |V|$ can solve the Hamiltonian Cycle problem in polynomial time.
Assuming $P \neq NP$, this is contradictory.

(Hamiltonian Cycle problem can be reduced to k cycle problem)

8.

Yes. (2%)

Explanation

Because if a directed graph has no cycles, we can use dynamic programming (e.g., Bellman-Ford style) along with the max-version of the RELAX operation to find the maximum distance from s to t .

Alternatively, we can also use a greedy approach based on the **topological ordering** of the vertices to find our solution.

Example Algorithm (8%)

1. Topological sort the vertices of G .

2. Initialize

- $dist[v] := \begin{cases} 0, & v = s \\ -\infty, & v \neq s \end{cases}$
- $pred[v] := \text{NIL}$

3. Relax edges in topo-order:

```
for each u in topological order:  
    for each v in outgoing(u):  
        if dist[v] < dist[u] + w(u,v):  
            dist[v] = dist[u] + w(u,v)  
            pred[v] = u
```

4. Outputs:

- Longest-path length = $dist[t]$
 - Setting it to infinity if t is unreachable.
- Longest-path:
 - follow $pred$ backward from t to s and then reverse (if t is reachable from s).